

Linear Programming Approximations

1 Generalization of Maximum Coverage Approximation

We will see that the same approximation we used last lecture to solve maximum coverage leads to a $(1 - \frac{1}{e})$ approximation for any non-negative monotone submodular function.

MAXIMUM COVERAGE: Given a universe U where each element has a value $v(u)$, subsets $S_1, \dots, S_m \subseteq U$, and an integer k , find a $T \subseteq \{1, \dots, m\}$ with $|T| \leq k$ maximizing $\sum_{u \in \bigcup_{j \in T} S_j} v(u)$.

Let function $f : 2^U \rightarrow \mathbb{R}$; in other words, $f : T \rightarrow \sum_{u \in \bigcup_{j \in T} S_j} v(u)$. Assume $f(T) \geq 0 \forall T$.

Monotonicity: $f(S) \geq f(T)$ when $S \supseteq T$.

Submodularity (diminishing returns): If $S \supseteq T$, then $f(T \cup \{x\}) - f(T) \geq f(S \cup \{x\}) - f(S) \forall x$. More elegantly, $f(S \cap T) + f(S \cup T) \leq f(S) + f(T)$.

Recall key insight 2 from last lecture: by adding the best set S_j for $j \in T^*$ to T_t , we increase the value by at least $\frac{1}{k}(OPT - V_t)$ because $|T^*| \leq k$.

Lemma 1. *If f is non-negative, submodular, and monotone, then this insight holds.*

Proof. Let $S = T^* \setminus T_t$; $S = \{u_1, \dots, u_r\}$ with $r \leq k$. Define $S_j := \{u_1, \dots, u_j\}$ and recall $f(T) = \sum_{u \in \bigcup_{j \in T} S_j} v(u)$. We are interested in:

$$f(T^* \cup T_t) - f(T_t)$$

$$= f(S \cup T_t) - f(T_t)$$

We use a telescoping sum:

$$\begin{aligned} &= \sum_{j=0}^{r-1} f(S_{j+1} \cup T_t) - f(S_j \cup T_t) \\ &= \sum_{j=0}^{r-1} f(S_j \cup T_t \cup \{u_{j+1}\}) - f(S_j \cup T_t) \end{aligned}$$

Applying submodularity:

$$\leq \sum_{j=0}^{r-1} f(T_t \cup \{u_{j+1}\}) - f(T_t)$$

We use this fact to prove the previous insight:

$$\begin{aligned}
\max_{j=1 \dots r} f(T_t \cup \{u_j\}) - f(T_t) &\geq \frac{1}{r} \sum_{j=1}^r f(T_t \cup \{u_j\}) - f(T_t) \\
&\geq \frac{1}{r} (f(T^* \cup T_t) - f(T_t))
\end{aligned}$$

And because $r \leq k$:

$$\geq \frac{1}{k} (f(T^* \cup T_t) - f(T_t))$$

□

Theorem 2. *Let f be a non-negative, monotone, submodular function. Then, the greedy algorithm which for k iterations always adds the element giving largest increase in f is a $1 - \frac{1}{e}$ approximation (Nemhauser, Halsey, Fisher).*

In the greedy algorithms unit, we proved that if f is **modular** (i.e., linear), maximizing $f(S)$ subject to S being an independent set in a given matroid is solved optimally by the greedy algorithm. The implications of Theorem 2 are that if f is non-negative, monotone, and submodular, the greedy algorithm is a $1 - \frac{1}{e}$ approximation on the k -uniform matroid (where any set of $\leq k$ elements is independent).

Theorem 3. *The greedy algorithm for a non-negative, monotone, submodular function on an arbitrary matroid is a $\frac{1}{2}$ -approximation (Nemhauser).*

Theorem 4. *There is a “continuous-greedy” (think gradient descent) algorithm which achieves a $1 - \frac{1}{e}$ approximation for any non-negative, monotone, submodular function on any matroid (Vondrak).*

2 Intro to Linear Programming Approximations

Many optimization problems can be encoded naturally as integer linear programs (ILPs). Solving ILPs is **NP**-hard because many problems are naturally encoded. The generic technique for achieving bounds on OPT:

1. Create the ILP encoding the problem.
2. “Relax” the ILP to an LP by removing integrality constraints.
3. Now the LP can be solved in polynomial time to give a fractional solution \vec{x} .
4. Round \vec{x} to an integer solution \hat{x} and ensure it is a valid solution (use \vec{x} as guidance for computing \hat{x}).
5. Prove \hat{x} is not too much worse than \vec{x} .

Important bounds:

$$OPT_{LP} \leq OPT_{ILP} \text{ for minimization}$$

$$OPT_{LP} \geq OPT_{ILP} \text{ for maximization}$$

Corollaries:

If $cost(\hat{x}) \leq \alpha * cost(\vec{x})$, then \hat{x} is an α -approximation for minimization.

If $cost(\hat{x}) \geq \alpha * cost(\vec{x})$, then \hat{x} is an α -approximation for maximization.

The **integrality gap** of an LP is, over all instances:

$$\min \frac{OPT_{ILP}}{OPT_{LP}} \text{ for minimization}$$

$$\max \frac{OPT_{ILP}}{OPT_{LP}} \text{ for maximization}$$

If analysis only uses LP bound, it can prove no better approximation guarantees than the integrality gap.

3 Weighted Vertex Cover

3.1 Building an LP

WEIGHTED VERTEX COVER: Given a graph $\mathcal{G} = (V, E)$ with vertex costs c_v , find a vertex cover S of minimum total cost $\sum_{u \in S} c_u$.

We write an ILP encoding the problem:

Maximize

$$\sum_v c_v x_v$$

subject to

$$\begin{cases} x_u + x_v \geq 1 & \forall e = (u, v) \\ x_v \geq 0 & \forall v \\ x_v \in \{0, 1\} & \forall v \end{cases}$$

We relax the ILP to an LP by removing the $x_v \in \{0, 1\}$ constraint. Solving this LP gives a fractional solution \vec{x} .

3.2 Rounding Algorithm

We include in S all nodes v with $x_v \geq \frac{1}{2}$.

(1). This is a vertex cover: because of the first constraint, at least one endpoint of e has $x_v \geq \frac{1}{2} \forall e$. Such a v is in S .

(2). Cost analysis:

$$\begin{aligned} cost(\hat{x}) &= \sum_{v \in S} c_v = \sum_{v: x_v \geq \frac{1}{2}} c_v \\ &\leq \sum_{v: x_v \geq \frac{1}{2}} (2x_v) c_v \end{aligned}$$

$$\begin{aligned}
&\leq \sum_v (2x_v) c_v \\
&= 2 * cost(\vec{x})
\end{aligned}$$

So this is a 2-approximation.

3.3 Integrality Gap

The worst case scenario is a complete graph with all $c_v = 1$. Then, $LP - OPT \leq \frac{n}{2}$ and $ILP - OPT = n - 1$. This ratio is about 2, so we know that the rounding can't be better than a 2-approximation.